

RICHARD OMOROTIONMWAN

Software Engineering | Data Engineering | AI & Cloud Platforms

Oakland, CA | iheoma.richard@gmail.com | +1-628-468-8729 | linkedin.com/in/richardomor | github.com/iheomach

PROFILE

Software and data engineer with production experience building full-stack applications, agentic AI systems, and end-to-end ETL pipelines across cloud platforms. Delivered data and infrastructure solutions at Meta, an investment management firm with \$160B+ AUM, and a large-scale energy company, spanning real-time analytics, Terraform-automated provisioning, and LLM-powered automation. Independently shipped Ogwu, an agentic telehealth platform with a multi-tool LangGraph agent, event-driven AWS pipeline, and an LLM-as-a-judge eval suite. Fluent in Python, SQL, PySpark, dbt, and Databricks, and comfortable translating messy data and ambiguous business problems into robust, repeatable systems.

TECHNICAL SKILLS

Languages: Python (Pandas, PySpark), SQL, JavaScript, TypeScript, Java, C++, Bash, PHP

Data Engineering & ETL: dbt, Apache Spark, PySpark, Databricks (Unity Catalog), Azure Data Factory, Azure Synapse, data modeling, schema design, data validation, warehouse architecture

AI & Agentic Systems: LangGraph, OpenAI API, Claude API, MCP, agent orchestration, tool-calling, RAG pipelines, pgvector

Cloud & Infrastructure: AWS (Lambda, SQS, EventBridge, S3, Textract, CloudWatch), Azure (ADF, Databricks, Synapse, ADLS), Terraform, Docker, GitHub Actions, Azure DevOps, CI/CD

Web & APIs: React, React Native, Next.js, Node.js/Express, REST APIs, webhooks, OAuth, event-driven architecture

Databases & BI: PostgreSQL, MySQL, Microsoft SQL Server, Supabase, Power BI, DAX, Power Query

EXPERIENCE

Infrastructure Support Engineer | Meta (contract via Crystal Equation)

Oct 2025 – Present

San Francisco, CA

- Support and maintain Scuba, Meta's real-time analytics platform, across 1,000+ production data pipelines, monitoring ingestion health at every stage, diagnosing ETL failures in distributed infrastructure, and driving remediation with backend and production engineers.
- Built **Claude**-powered diagnostic agents on Meta's internal agentic framework using **LLM tool-calling** and **retrieval**, automating failure investigation across pipelines and directly reducing mean time to resolution.
- Owned **storage capacity governance** for Scuba tables, driving allocation decisions across engineering and finance stakeholders while balancing reliability against cost.
- Built **React.js** features for internal operational tooling with input-validation logic that prevents malformed table storage requests, and added unit and **CI/CD** test coverage to verify regression fixes.

Data & AI Platform Engineer Intern | Alberta Investment Management Corporation (AIMCo)

Sep 2024 – Oct 2025

Edmonton, AB

- Designed and delivered **end-to-end data infrastructure on Azure** for a \$160B+ AUM investment firm: Terraform-automated Databricks Unity Catalog provisioning, ADLS External Locations, Storage Credentials, and service principal role assignments, eliminating manual release steps.
- Built an **Azure DevOps CI/CD pipeline** to automate dbt job deployment, improving pipeline delivery reliability and reducing release friction.
- Monitored ADF and Databricks ETL jobs daily for platform failures, led incident bridge calls, performed root-cause analysis, and drove remediation to maintain SLA compliance across investment data pipelines.
- Built an **automated analytics dashboard** integrating Databricks, SQL, Power BI, and Microsoft Copilot AI, cutting reporting time by 93%; presented the solution to business stakeholders across IT and finance.

Data Analyst Co-op | Suncor Energy

May 2023 – Aug 2024

Calgary, AB

- Built an **end-to-end automated financial reporting pipeline** using Power Automate, Azure Blob Storage, Databricks, and REST APIs with PySpark ETL transformations, cutting manual reporting effort by 50%.
- Developed an **OOP Python application** to automate risk hierarchy generation from parent-child relationships in a SQL database, replacing a manual workflow and feeding a live Power BI supply chain dashboard.
- Designed **DAX and Power Query scripts** for efficient querying in Power BI, producing operational reports tracking key supply chain and financial metrics for engineering and finance stakeholders.
- Designed and delivered **Power BI finance reporting dashboards**, translating raw financial data into clear visualizations and actionable summaries for business decision-making.

PROJECTS

Ogwu | LangGraph, Node.js, React Native, OpenAI, AWS, Supabase, pgvector

github.com/iheomach/ogwu

- Built an **8-tool LangGraph agentic platform** integrating hospital search by GPS, Google Calendar booking, drug interaction checks, emergency escalation, and patient history retrieval, demonstrating end-to-end multi-service agent orchestration.
- Designed an **event-driven AWS data pipeline** (S3, EventBridge, two Lambdas, SQS, DLQ, CloudWatch alarms) for document ingestion and processing, with full observability and least-privileged IAM roles per function.
- Built a **RAG pipeline** using AWS Textract for document extraction, overlapping chunking into pgvector, and cosine-similarity retrieval for context injection into downstream workflows.
- Implemented an **LLM-as-a-judge CI eval suite** (60 test cases, 85% pass threshold) triggered on every graph change via GitHub Actions, cutting iteration time by 40%; shipped a full React Native patient app and Next.js doctor dashboard with multilingual support.

EDUCATION

M.S. Computer Science, ML Specialization | Georgia Institute of Technology (OMSCS)

Starting Fall 2026

Part-time | Expected Dec 2029

B.Sc. Computer Science | University of Alberta

2020 – 2025

CERTIFICATIONS

Microsoft Azure AI Engineer Associate (AI-102) | Microsoft Azure Fundamentals (AZ-900)